

Regression_PenalizedModel

April 9, 2022

```
[1]: import pandas as pd

data = pd.read_csv('C:\\Users\\MSI_\\
↪Stealth\\Downloads\\BMEN415Project\\regression\\Volumetric_features.csv')

print(data.head())
```

	S.No	Left-Lateral-Ventricle	Left-Inf-Lat-Vent	\
0	1	22916.9	982.7	
1	2	22953.2	984.5	
2	3	23320.4	1062.1	
3	4	24360.0	1000.5	
4	5	25769.4	1124.4	

	Left-Cerebellum-White-Matter	Left-Cerebellum-Cortex	Left-Thalamus	\
0	15196.7	55796.4	6855.5	
1	15289.7	55778.6	6835.1	
2	15382.1	55551.2	7566.0	
3	14805.4	54041.8	8004.6	
4	16331.1	54108.6	6677.4	

	Left-Caudate	Left-Putamen	Left-Pallidum	3rd-Ventricle	...	\
0	2956.4	4240.7	2223.9	2034.4	...	
1	3064.2	4498.6	2354.1	1927.1	...	
2	3231.7	4456.2	1995.4	2064.7	...	
3	3137.3	4262.2	1983.4	2017.7	...	
4	2964.4	4204.6	2409.7	2251.8	...	

	rh_supramarginal_thickness	rh_frontalpole_thickness	\
0	2.408	2.629	
1	2.417	2.640	
2	2.374	2.601	
3	2.366	2.639	
4	2.381	2.555	

	rh_temporalpole_thickness	rh_transversetemporal_thickness	\
0	3.519	2.009	
1	3.488	2.111	

2	3.342	2.146
3	3.361	2.056
4	3.450	2.052

	rh_insula_thickness	rh_MeanThickness_thickness	BrainSegVolNotVent.2 \
0	2.825	2.33635	1093846
1	2.720	2.34202	1099876
2	2.684	2.31982	1097999
3	2.700	2.29215	1070117
4	2.574	2.30397	1075926

	eTIV.1	Age	dataset
0	1619602.965	85	1
1	1624755.130	85	1
2	1622609.518	86	1
3	1583854.236	87	1
4	1617375.362	89	1

[5 rows x 141 columns]

```
[2]: # Separate Target Variable and Predictor Variables
y=data['Age'].values
X=data[['Left-Lateral-Ventricle', 'Left-Inf-Lat-Vent',
        'Left-Cerebellum-White-Matter', 'Left-Cerebellum-Cortex',
        'Left-Thalamus', 'Left-Caudate', 'Left-Putamen',
        'Left-Pallidum', '3rd-Ventricle', '4th-Ventricle',
        'Brain-Stem', 'Left-Hippocampus', 'Left-Amygdala',
        'CSF', 'Left-Accumbens-area', 'Left-VentralDC',
        'Left-vessel', 'Left-choroid-plexus', 'Right-Lateral-Ventricle',
        'Right-Inf-Lat-Vent', 'Right-Cerebellum-White-Matter', 'Right-Cerebellum-Cortex',
        'Right-Thalamus', 'Right-Caudate', 'Right-Putamen',
        'Right-Pallidum', 'Right-Hippocampus', 'Right-Amygdala', 'Right-Accumbens-area',
        'Right-VentralDC', 'Right-vessel', 'Right-choroid-plexus',
        '5th-Ventricle', 'WM-hypointensities', 'Left-WM-hypointensities',
        'Right-WM-hypointensities', 'non-WM-hypointensities', 'Left-non-WM-hypointensities',
        'Right-non-WM-hypointensities', 'Optic-Chiasm', 'CC_Posterior', 'CC_Mid_Posterior',
        'CC_Central', 'CC_Mid_Anterior', 'CC_Anterior', 'BrainSegVol', 'BrainSegVolNotVent']].values
```

```
[3]: train = data[:4000]
test = data[4000:]
X_test = X[4000:]
```

```

y_test = y[4000:]
X_train = X[:4000]
y_train = y[:4000]

```

```

[4]: import numpy as np
from sklearn.linear_model import LogisticRegression
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from patsy import dmatrices, dmatrix
import matplotlib.pyplot as plt
import math

```

```

# Create a scaler object
sc = StandardScaler()

```

```

# Fit the scaler to the training data and transform
X_train_std = sc.fit_transform(X_train)

```

```

# Apply the scaler to the test data
X_test_std = sc.transform(X_test)


```

```

[5]: from sklearn.linear_model import LassoCV, ElasticNetCV, RidgeCV
reg_lasso = LassoCV(cv=10).fit(X_train_std, y_train)

EPSILON = 1e-4 # This is to avoid division by zero while taking the base 10
↳ logarithm
plt.figure()
plt.semilogx(reg_lasso.alphas_ + EPSILON, reg_lasso.mse_path_, ':')
plt.plot(reg_lasso.alphas_ + EPSILON, reg_lasso.mse_path_.mean(axis=-1), 'k',
         label='Average across the folds', linewidth=2)
plt.axvline(reg_lasso.alpha_ + EPSILON, linestyle='--', color='k',
            label=r'$\alpha$: CV estimate')
plt.legend()
plt.xlabel(r'$\alpha$')
plt.ylabel('Mean square error')
plt.title('Mean square error on each fold: coordinate descent ')
plt.axis('tight')
plt.show()
print ('Optimized alpha', reg_lasso.alphas_)

```



Optimized alpha [14.08755217 13.1380983 12.25263445 11.42684789 10.65671659
9.93848956

9.26866863	8.64399138	8.06141529	7.51810288	7.01140791	6.53886248
6.09816504	5.68716912	5.30387296	4.94640967	4.61303822	4.30213488
4.01218539	3.74177752	3.48959424	3.25440727	3.03507111	2.83051747
2.63975007	2.46183975	2.29591999	2.14118267	1.99687412	1.86229149
1.73677928	1.61972617	1.51056204	1.40875522	1.31380983	1.22526344
1.14268479	1.06567166	0.99384896	0.92686686	0.86439914	0.80614153
0.75181029	0.70114079	0.65388625	0.6098165	0.56871691	0.5303873
0.49464097	0.46130382	0.43021349	0.40121854	0.37417775	0.34895942
0.32544073	0.30350711	0.28305175	0.26397501	0.24618398	0.229592
0.21411827	0.19968741	0.18622915	0.17367793	0.16197262	0.1510562
0.14087552	0.13138098	0.12252634	0.11426848	0.10656717	0.0993849
0.09268669	0.08643991	0.08061415	0.07518103	0.07011408	0.06538862
0.06098165	0.05687169	0.05303873	0.0494641	0.04613038	0.04302135
0.04012185	0.03741778	0.03489594	0.03254407	0.03035071	0.02830517
0.0263975	0.0246184	0.0229592	0.02141183	0.01996874	0.01862291
0.01736779	0.01619726	0.01510562	0.01408755]		

```
[6]: from sklearn.linear_model import Lasso
reg = Lasso(alpha=0.01408755)
reg.fit(X_train_std, y_train)
```

```
Lasso(alpha=0.01408755)
print('R squared training set', round(reg.score(X_train_std, y_train)*100, 2))
print('R squared test set', round(reg.score(X_test_std, y_test)*100, 2))
```

R squared training set 81.59

R squared test set -89.09

```
[8]: from sklearn.metrics import mean_squared_error
      predictions = reg.predict(X_test_std)
      total_cases = len(y_test) # size of validation set
```

```
avg = 0.0
SSres = 0.0
SStot = 0.0
```

```
for i in range(total_cases):
    value = y_test[i]
    predict = predictions[i]
    avg = (avg + value)/2
    SSres = SSres + (value - predict)**2
    SStot = SStot + (value - avg)**2
    #print(value, '----- ', predict)
```

```
Rsquared_cal = 1 - SStot/SSres
```

```
print('R squared: ', Rsquared_cal)
mse = mean_squared_error(y_test, reg.predict(X_test_std))
rmse = math.sqrt(mse)
print("The rooted mean squared error (MSE) on test set:{:.4f}", rmse)
```

R squared: 0.9084155078317926

The rooted mean squared error (MSE) on test set:{:.4f} 9.545207751552638

```
[ ]:
```